

PyTorch Based Implementation of Semi-Model-Based Reinforcement Learning with Uncertainty

Simon Lund, Sophia Sigethy, Georg Staber, and Malte Wilhelm

28.09.2021

Ludwig Maximilian University of Munich

Table of Contents

1. Introduction
2. Concept
3. Experiments & Results
4. Conclusion



Introduction

Motivation

Real life data is often **noisy, scarce and unreliable**

Motivation

Real life data is often **noisy, scarce and unreliable**

Prediction quality of neural networks is highly dependent on **quality of training data**

Motivation

Real life data is often **noisy, scarce and unreliable**

Prediction quality of neural networks is highly dependent on **quality of training data**


Model becomes **uncertain** regarding $y = f(x)$

Motivation

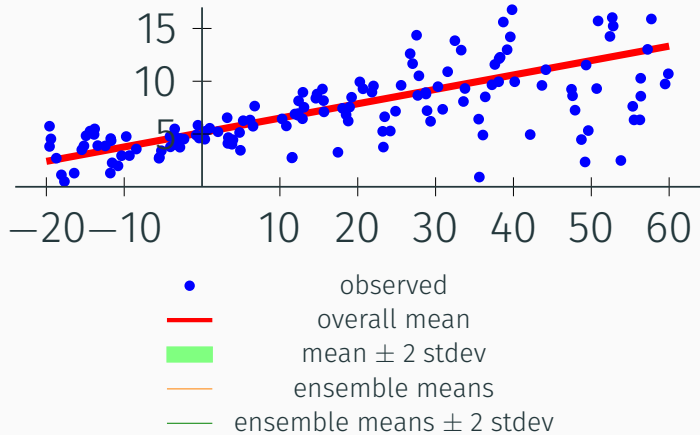
Real life data is often **noisy, scarce and unreliable**

Prediction quality of neural networks is highly dependent on **quality of training data**

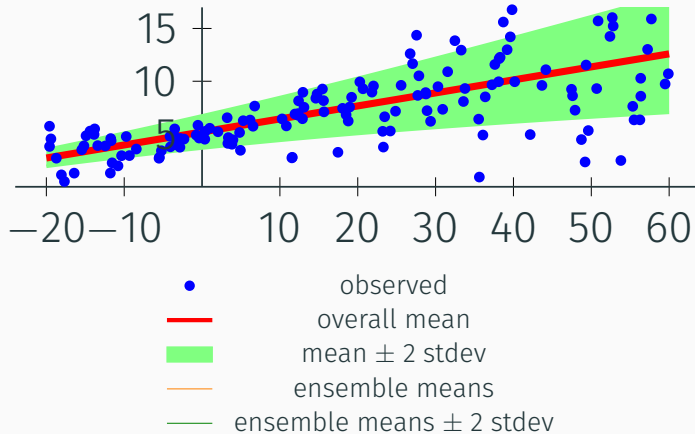
Model becomes **uncertain regarding $y = f(x)$**




No Uncertainty (linear regression)



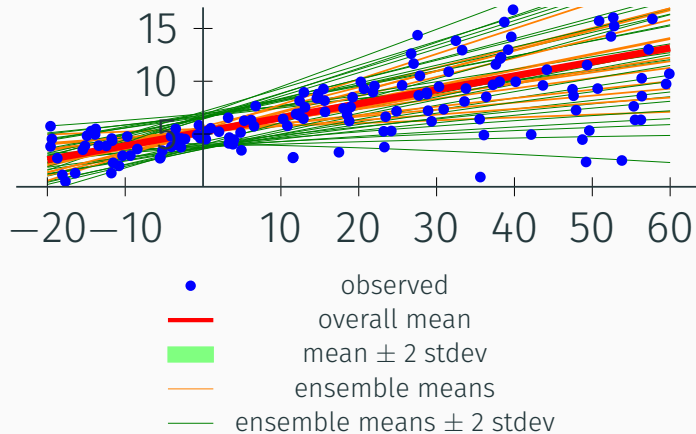
Aleatoric Uncertainty



Epistemic Uncertainty



Aleatoric and Epistemic Uncertainty



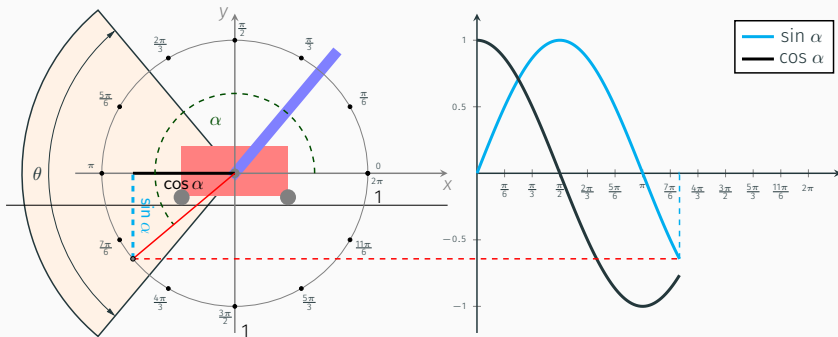
Concept

1. Sample data from noisy Cartpole Environment


1. Sample data from noisy Cartpole Environment
2. Learn neural network for one-step dynamics

1. Sample data from noisy Cartpole Environment
2. Learn neural network for one-step dynamics
3. Train RL policy against dynamics model

Data Sampling



Neural Network Architecture

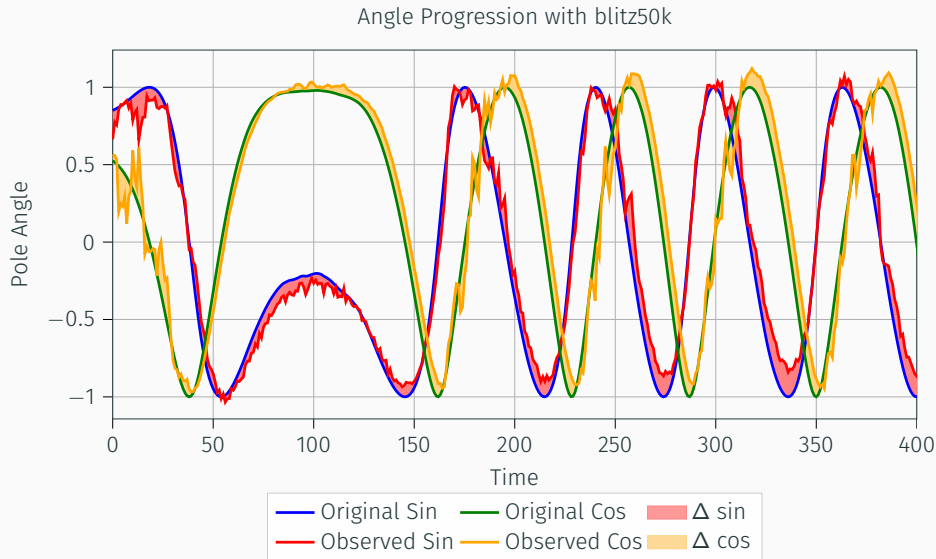


Experiments & Results


BNN predictions: 1k time series





BNN predictions: 50k time series



BNN Uncertainty: 75k time series



BNN Uncertainty: 75k time series



With Uncertainty

Without uncertainty

LIVE DEMO

Conclusion

1. RL agent detects noisy sector and avoids it

1. RL agent detects noisy sector and avoids it
2. RL agent is able to "solve" the environment even with uncertainty




Take Aways


1. RL agent detects noisy sector and avoids it
2. RL agent is able to "solve" the environment even with uncertainty
3. Use PyTorch ;-)

Questions?

References

-  P. Sountsov, C. Suter, J. Burnim and J. V. Dillon. *Regression with Probabilistic Layers in TensorFlow Probability* — *The TensorFlow Blog*. 12th Mar. 2019. URL: <https://blog.tensorflow.org/2019/03/regression-with-probabilistic-layers-in.html> (visited on 2nd Sept. 2021).
-  A. Raffin, A. Hill, M. Ernestus, A. Gleave, A. Kanervisto and N. Dormann. *Stable Baselines3*. <https://github.com/DLR-RM/stable-baselines3>. 2019.

-  Â. Lovatto. *angelolovatto/gym-cartpole-swingup: A simple, continuous-control environment for OpenAI Gym*. 7th July 2019. URL: <https://github.com/angelolovatto/gym-cartpole-swingup> (visited on 2nd Sept. 2021).
-  G. Brockman, V. Cheung, L. Pettersson, J. Schneider, J. Schulman, J. Tang and W. Zaremba. "OpenAI Gym". In: (2016). eprint: *arXiv:1606.01540*. URL: <https://gym.openai.com/>.
-  C. Blundell, J. Cornebise, K. Kavukcuoglu and D. Wierstra. *Weight Uncertainty in Neural Networks*. 2015. arXiv: *1505.05424 [stat.ML]*.

-  A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Kopf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai and S. Chintala. “PyTorch: An Imperative Style, High-Performance Deep Learning Library”. In: *Advances in Neural Information Processing Systems* 32. Ed. by H. Wallach, H. Larochelle, A. Beygelzimer, F. d’Alché-Buc, E. Fox and R. Garnett. Curran Associates, Inc., 2019, pp. 8024–8035. URL: <http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf>.

Backup Slides

Creating a variational regressor class

```
1  @variational_estimator
2  class BayesianRegressor(nn.Module):
3      def __init__(self, input_dim, output_dim):
4          super().__init__()
5          self.blinear = BayesianLinear(input_dim, output_dim)
6
7      def forward(self, x):
8          return self.blinear(x)
```

Defining a confidence interval evaluating function

```
1  def evaluate_regression(regressor, x, y, samples=25, std_multiplier=2):
2      preds = [regressor(x) for i in range(samples)]
3      preds = torch.stack(preds)
4      means = preds.mean(axis=0)
5      stds = preds.std(axis=0)
6      ci_upper = means + (std_multiplier * stds)
7      ci_lower = means - (std_multiplier * stds)
8      ic_acc = (ci_lower <= y) * (ci_upper >= y)
9      ic_acc = ic_acc.float().mean()
10     return ic_acc, (ci_upper >= y).float().mean(), (ci_lower <= y).float().mean()
```

Creating our regressor and loading data

```
1 optimizer = optim.Adam(regressor.parameters(), lr=0.01)
2 criterion = torch.nn.MSELoss()
3 complexity_cost_weight = 1. / x_train.shape[0]
```

Main training and evaluating loop

```
1 losses = []
2 for epoch in tqdm(range(100)):
3     new_epoch = True
4     for i, (datapoints, labels) in enumerate(dataloader_train):
5         optimizer.zero_grad()
6
7         loss = regressor.sample_elbo(
8             inputs=datapoints,
9             labels=labels,
10            criterion=criterion,
11            sample_nbr=3,
12            complexity_cost_weight=complexity_cost_weight
13        )
14
15        loss.backward()
16        optimizer.step()
```